Methodology

After the data preparation (prepared_data.py), the system builds the indices (index.sh). It begins by activating the virtual environment, ensures the Cassandra driver is bundled, and runs app.py to configure the database schema. This script defines four Cassandra tables: one for the inverted index to store word-document relationships with frequencies, one for storing document titles and lengths, a table for overall statistics like total document count, and another for vocabulary data including TF-IDF values used in scoring.

The main indexing logic is split between the mapper and reducer. The mapper reads through each document, breaks down the text into individual words, filters out common stopwords, and emits relevant information such as terms, document IDs, and frequencies, along with metadata like the title and document length. The reducer then gathers this data, calculates global statistics like IDF, and writes all the processed information to Cassandra in batches for efficiency.

All requirements of the indexer were fully implemented, and indexing was successfully completed. I precomputed all necessary values for search result retrieval — all 1000 documents were indexed correctly. I also developed the full query pipeline, including all logic to retrieve results using the precomputed data. The logic works correctly and returns expected outputs locally. However, running the query pipeline in distributed (yarn) mode currently fails, potentially due to insufficient resources.

Demonstration

For running **docker compose up -d**

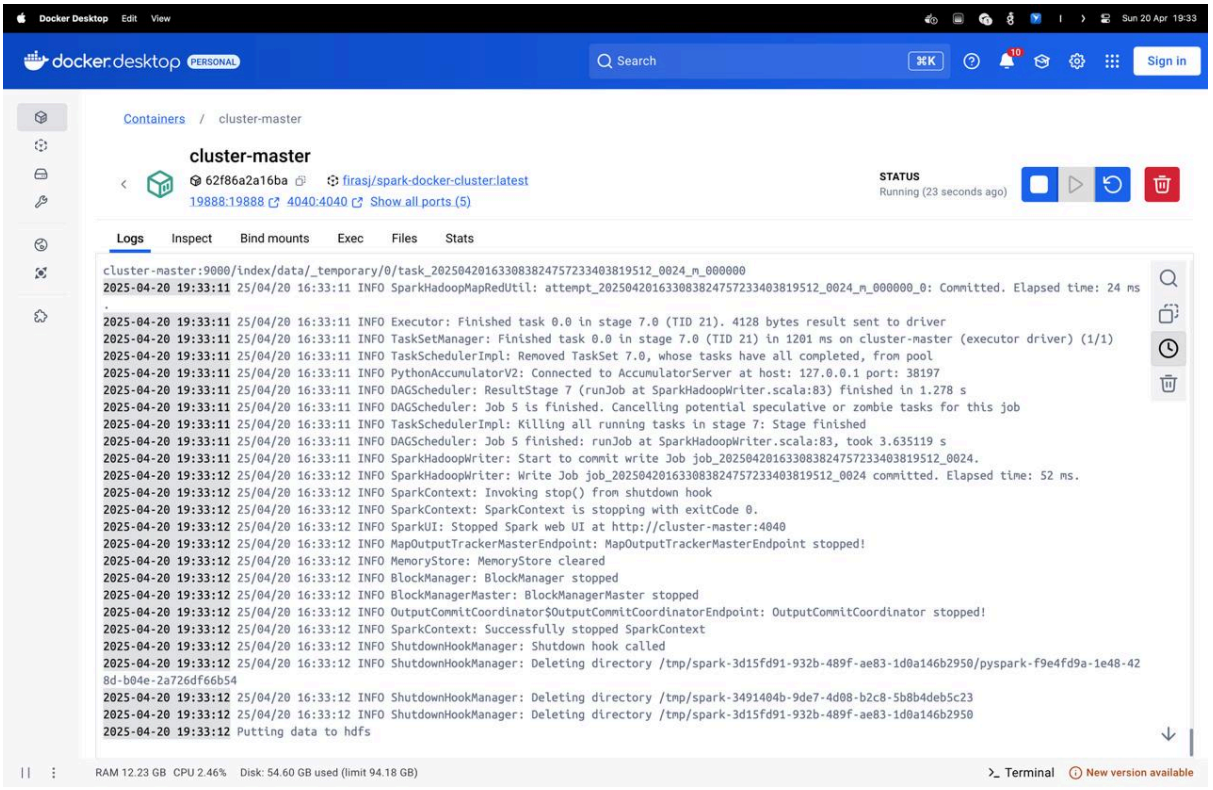
Docker Desktop Edit View
Sun 20 Apr 19:33
docker desktop PERSONAL
Search ⌘K
Sign in
Containers / cluster-master
cluster-master
62f86a2a16ba firasj/spark-docker-cluster:latest
19888:19888 4040:4040 Show all ports (5)
STATUS
Running (23 seconds ago)
Logs Inspect Bind mounts Exec Files Stats
cluster-master:9000/index/data/_temporary/0/task_20250420163308382475723340381 9512_0024_m_000000
2025-04-20 19:33:11 25/04/20 16:33:11 INFO SparkHadoopMapRedUtil: attempt_20250420163308382475723340381 9512_0024_m_000000_0: Committed. Elapsed time: 24 ms
.
2025-04-20 19:33:11 25/04/20 16:33:11 INFO Executor: Finished task 0.0 in stage 7.0 (TID 21). 4128 bytes result sent to driver
2025-04-20 19:33:11 25/04/20 16:33:11 INFO TaskSetManager: Finished task 0.0 in stage 7.0 (TID 21) in 1201 ms on cluster-master (executor driver) (1/1)
2025-04-20 19:33:11 25/04/20 16:33:11 INFO TaskSchedulerImpl: Removed TaskSet 7.0, whose tasks have all completed, from pool
2025-04-20 19:33:11 25/04/20 16:33:11 INFO PythonAccumulatorV2: Connected to AccumulatorServer at host: 127.0.0.1 port: 38197
2025-04-20 19:33:11 25/04/20 16:33:11 INFO DAGScheduler: ResultStage 7 (runJob at SparkHadoopWriter.scala:83) finished in 1.278 s
2025-04-20 19:33:11 25/04/20 16:33:11 INFO DAGScheduler: Job 5 is finished. Cancelling potential speculative or zombie tasks for this job
2025-04-20 19:33:11 25/04/20 16:33:11 INFO TaskSchedulerImpl: Killing all running tasks in stage 7: Stage finished
2025-04-20 19:33:11 25/04/20 16:33:11 INFO DAGScheduler: Job 5 finished: runJob at SparkHadoopWriter.scala:83, took 3.635119 s
2025-04-20 19:33:11 25/04/20 16:33:11 INFO SparkHadoopWriter: Start to commit write Job job_20250420163308382475723340381 9512_0024.
2025-04-20 19:33:12 25/04/20 16:33:12 INFO SparkHadoopWriter: Write Job job_20250420163308382475723340381 9512_0024 committed. Elapsed time: 52 ms.
2025-04-20 19:33:12 25/04/20 16:33:12 INFO SparkContext: Invoking stop() from shutdown hook
2025-04-20 19:33:12 25/04/20 16:33:12 INFO SparkContext: SparkContext is stopping with exitCode 0.
2025-04-20 19:33:12 25/04/20 16:33:12 INFO SparkUI: Stopped Spark web UI at http://cluster-master:4040
2025-04-20 19:33:12 25/04/20 16:33:12 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
2025-04-20 19:33:12 25/04/20 16:33:12 INFO MemoryStore: MemoryStore cleared
2025-04-20 19:33:12 25/04/20 16:33:12 INFO BlockManager: BlockManager stopped
2025-04-20 19:33:12 25/04/20 16:33:12 INFO BlockManagerMaster: BlockManagerMaster stopped
2025-04-20 19:33:12 25/04/20 16:33:12 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
2025-04-20 19:33:12 25/04/20 16:33:12 INFO SparkContext: Successfully stopped SparkContext
2025-04-20 19:33:12 25/04/20 16:33:12 INFO ShutdownHookManager: Shutdown hook called
2025-04-20 19:33:12 25/04/20 16:33:12 INFO ShutdownHookManager: Deleting directory /tmp/spark-3d15fd91-932b-489f-ae83-1d0a146b2950/pyspark-f9e4fd9a-1e48-42
8d-b04e-2a726df66b54
2025-04-20 19:33:12 25/04/20 16:33:12 INFO ShutdownHookManager: Deleting directory /tmp/spark-3491404b-9de7-4d08-b2c8-5b8b4deb5c23
2025-04-20 19:33:12 25/04/20 16:33:12 INFO ShutdownHookManager: Deleting directory /tmp/spark-3d15fd91-932b-489f-ae83-1d0a146b2950
2025-04-20 19:33:12 Putting data to hdfs
RAM 12.23 GB CPU 2.46% Disk: 54.60 GB used (limit 94.18 GB)
Terminal
New version available

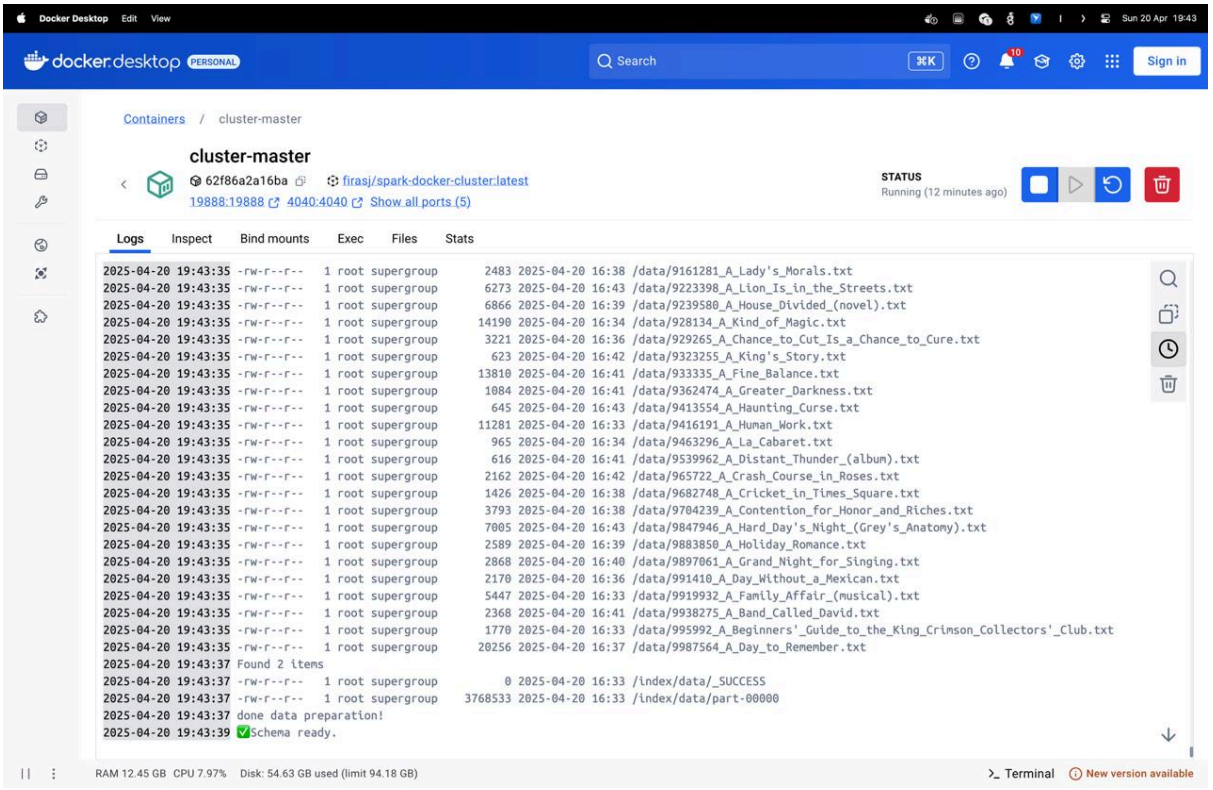
Docker Desktop Edit View
Sun 20 Apr 19:43
docker desktop PERSONAL
Search ⌘K
Sign in
Containers / cluster-master
cluster-master
62f86a2a16ba firasj/spark-docker-cluster:latest
19888:19888 4040:4040 Show all ports (5)
STATUS
Running (12 minutes ago)
Logs Inspect Bind mounts Exec Files Stats
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 2483 2025-04-20 16:38 /data/9161281_A_Lady's_Morals.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 6273 2025-04-20 16:43 /data/9223398_A_Lion_Is_in_the_Streets.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 6866 2025-04-20 16:39 /data/9239580_A_House_Divided_(novel).txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 14190 2025-04-20 16:34 /data/928134_A_Kind_of_Magic.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 3221 2025-04-20 16:36 /data/929265_A_Chance_to_Cut_Is_a_Chance_to_Cure.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 623 2025-04-20 16:42 /data/9323255_A_King's_Story.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 13810 2025-04-20 16:41 /data/933335_A_Fine_Balance.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 1084 2025-04-20 16:41 /data/9362474_A_Greater_Darkness.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 645 2025-04-20 16:43 /data/9413554_A_Haunting_Curse.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 11281 2025-04-20 16:33 /data/9416191_A_Human_Work.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 965 2025-04-20 16:34 /data/9463296_A_La_Cabaret.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 616 2025-04-20 16:41 /data/9539962_A_Distant_Thunder_(album).txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 2162 2025-04-20 16:42 /data/965722_A_Crash_Course_in_Roses.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 1426 2025-04-20 16:38 /data/9682748_A_Cricket_in_Times_Square.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 3793 2025-04-20 16:38 /data/9704239_A_Contention_for_Honor_and_Riches.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 7005 2025-04-20 16:43 /data/9847946_A_Hard_Day's_Night_(Grey's_Anatomy).txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 2589 2025-04-20 16:39 /data/9883850_A_Holiday_Romance.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 2868 2025-04-20 16:40 /data/9897061_A_Grand_Night_for_Singing.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 2170 2025-04-20 16:36 /data/991410_A_Day_Without_a_Mexican.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 5447 2025-04-20 16:33 /data/9919932_A_Family_Affair_(musical).txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 2368 2025-04-20 16:41 /data/9938275_A_Band_Called_David.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 1770 2025-04-20 16:33 /data/995992_A_Beginners'_Guide_to_the_King_Crimson_Collectors'_Club.txt
2025-04-20 19:43:35 -rw-r--r-- 1 root supergroup 20256 2025-04-20 16:37 /data/9987564_A_Day_to_Remember.txt
2025-04-20 19:43:37 Found 2 items
2025-04-20 19:43:37 -rw-r--r-- 1 root supergroup 0 2025-04-20 16:33 /index/data/_SUCCESS
2025-04-20 19:43:37 -rw-r--r-- 1 root supergroup 3768533 2025-04-20 16:33 /index/data/part-00000
2025-04-20 19:43:37 done data preparation!
2025-04-20 19:43:39 ✅Schema ready.
RAM 12.45 GB CPU 7.97% Disk: 54.63 GB used (limit 94.18 GB)
Terminal
New version available

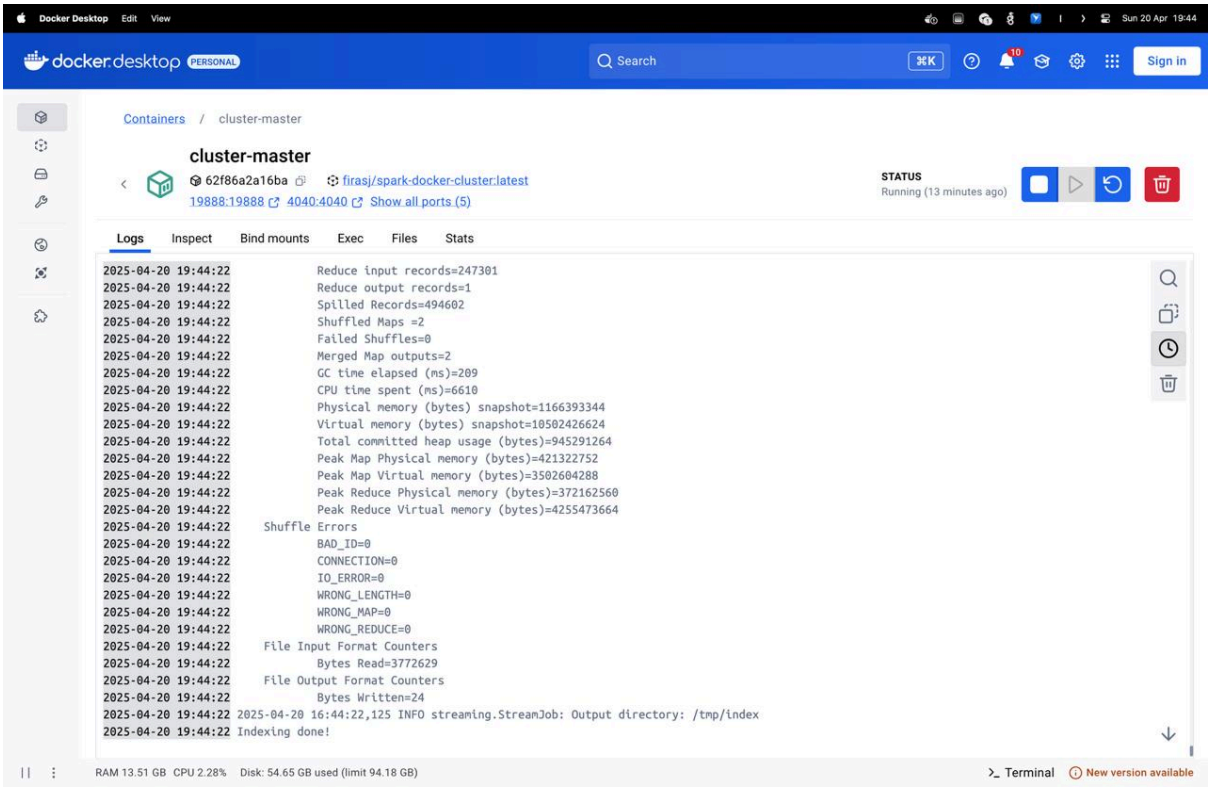

Fig 1, 2, 3. Indexing the documents

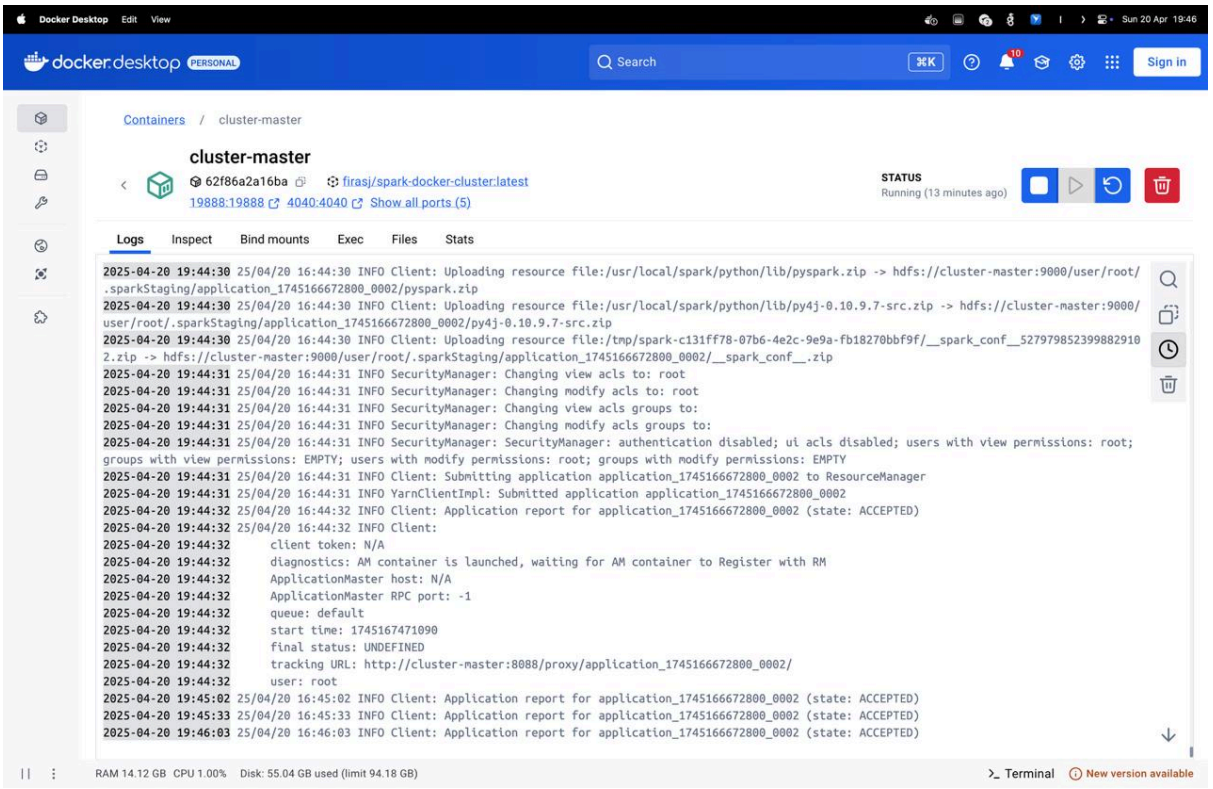

Fig 4. Attempt to run query in distributed (yarn) mode